

Social Media & Text Analysis

lecture 4 - natural language processing (part 2):
tokenization and normalization



Instructor: Wei Xu


Website: socialmedia-class.org

BAD LANGUAGE!

...on the INTERNET!!



Jacob **EISENSTEIN**
GEORGIA Institute of **TECH**nology



Boom! Ya ur
website suxx bro

...dats why pluto is pluto
it can neva be a star



michelle obama great.
job. and. whit all my.
respect she. look. great.
congrats. to. her.



I now h v an iphone

What can we do about it?
*Why don't they just write **NORMALLY??***
*Can our software ever **ADAPT???***

How does language go bad?

Illiteracy? No.

(Tagliamonte and Denis 2008;
Drouin and Davis 2009)




rob delaney @robdelaney

1 Jun

Great. Now a bunch of illiterate teens claim to be "powning" me with their insults. Heads up jerks my wife & children love me & are proud of

Expand Reply Classic RT Retweet Favorite More

Length limits? (probably not)

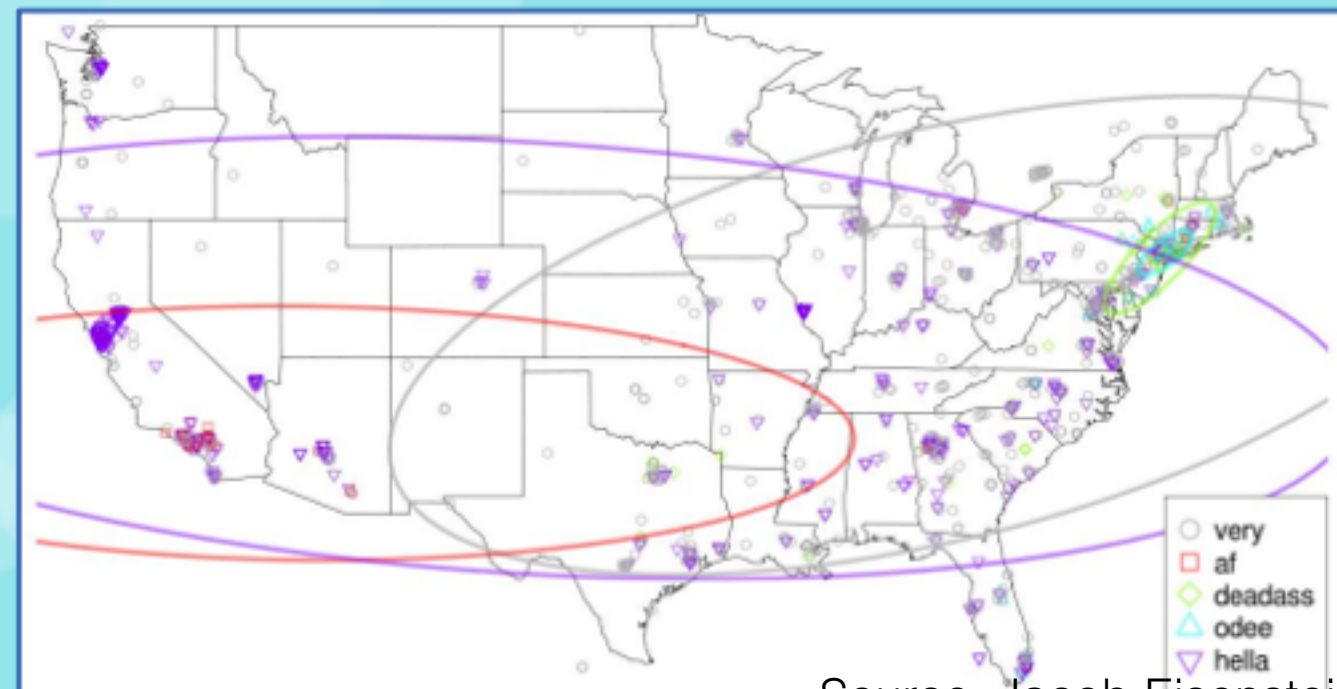


Hardware input constraints? (Gouws et al 2011)



Social variables

- Non-standard language does *identity work*, signaling authenticity, solidarity, etc.
- Social variation is usually inhibited in written language, but social media is less regulated than other written genres.

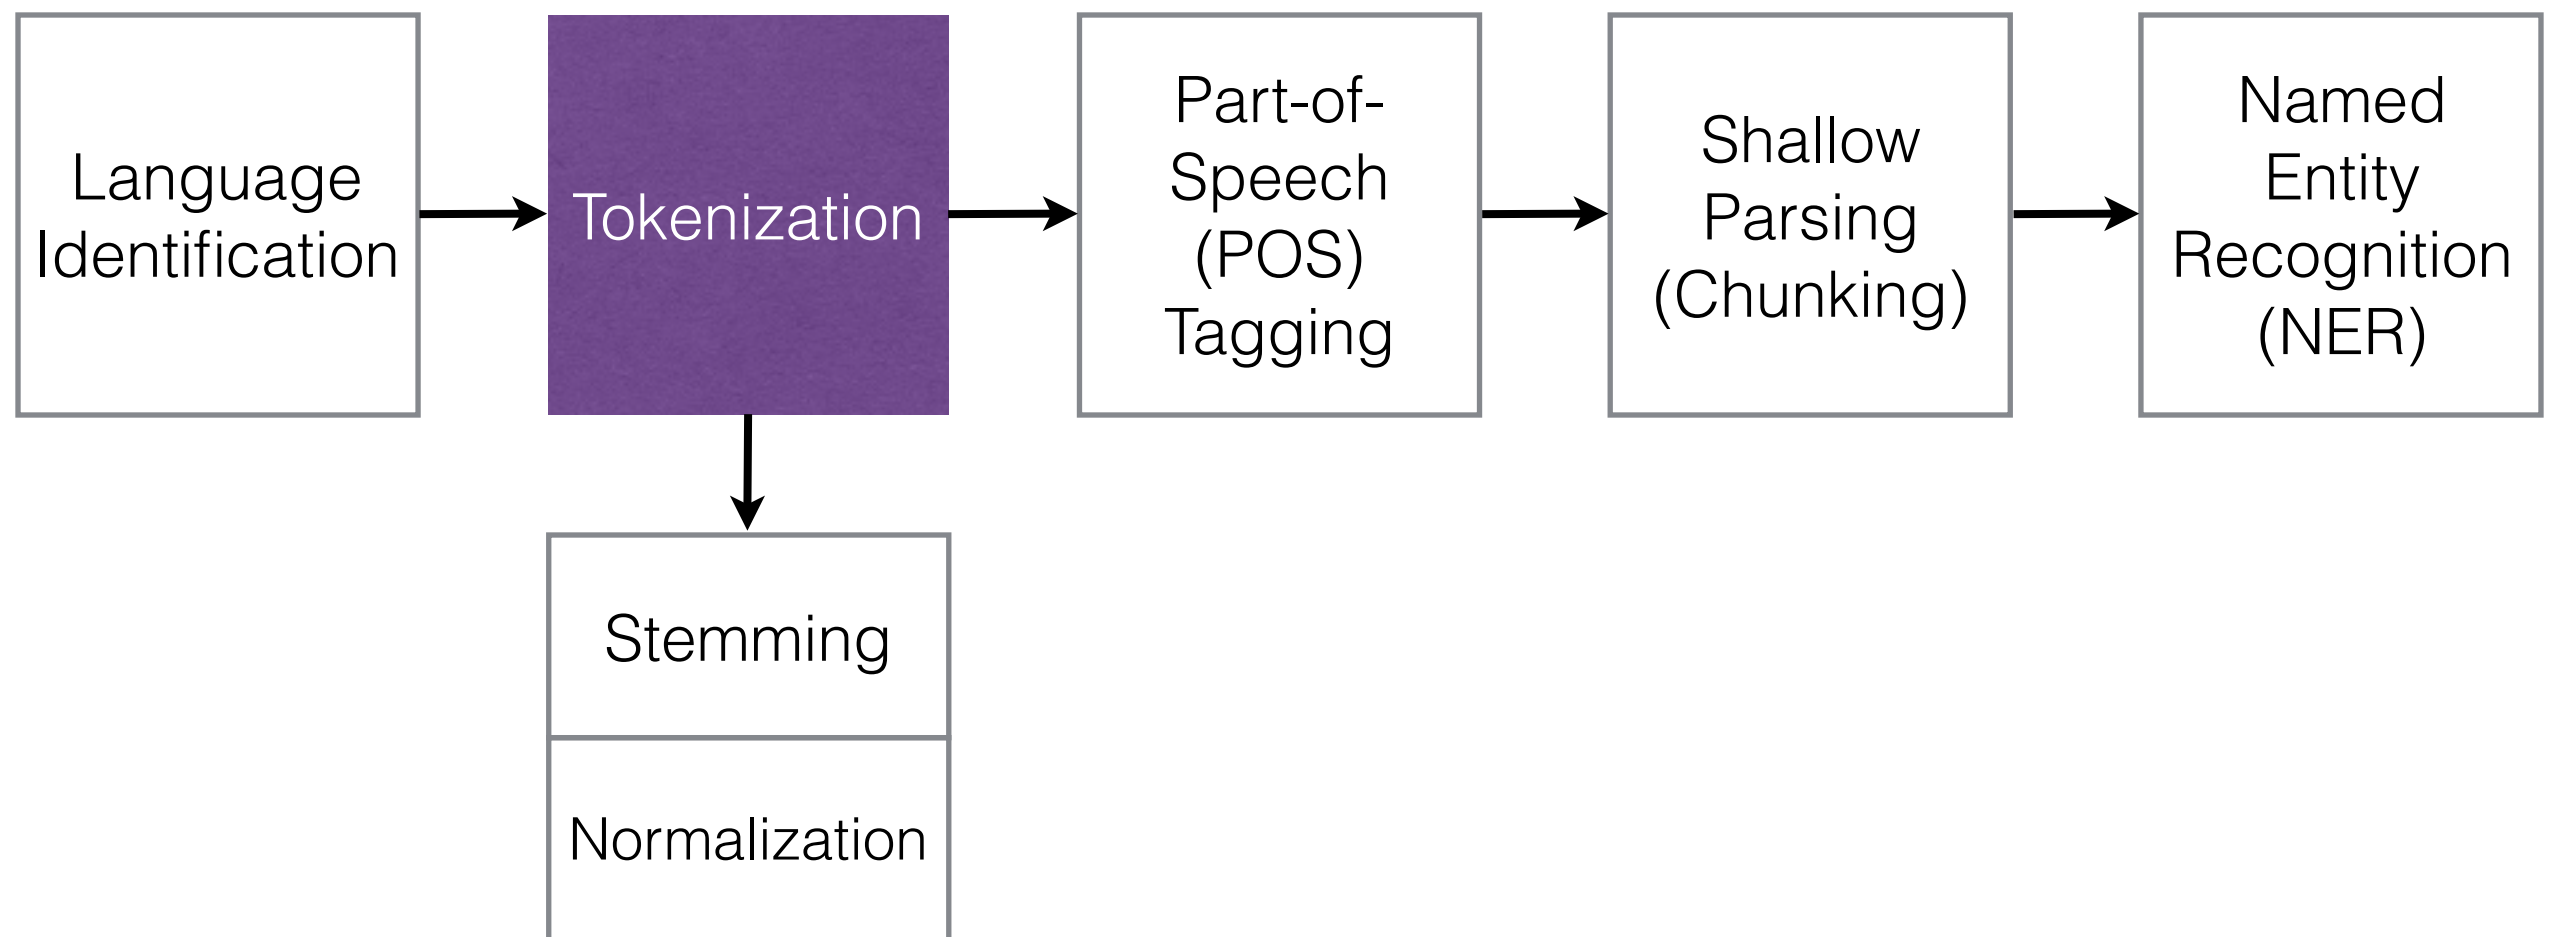


Source: Jacob Eisenstein

Why is Social Media Text “Bad”?

- Lack of literacy? **no** [Drouin and Davis, 2009]
- Length restrictions? **not primarily** [Eisenstein, 2013]
- Text input method? **to some degree, yes** [Gouws et al., 2011]
- Pragmatics (mimicking prosodic effects etc. in speech)? **yeeeees** [Eisenstein, 2013]
- Social variables/markers of social identity? **blood oath!** [Eisenstein, 2013]

NLP Pipeline



Tokenization

- breaks up the string into words and punctuation
- need to handle:
 - abbreviations (“jr.”), number (“5,000”) ...

```
seas479:training weixu$ ./penn-treebank-tokenizer.perl
Tokenizer v3
Language: en
```

```
Ms. Hilton last year called Mr. Rothschild "the love of my life."
```

```
Ms. Hilton last year called Mr. Rothschild " the love of my life . "
```

← **input**
← **output**

Tokenization


- for Twitter, additionally need to handle:
 - emoticons, urls, #hashtags, @mentions ...


```
>>> import twokenize
>>> input = "Clowns are pretty gross tho O.o (I'm afraid of clowns :p) ask.fm/a/cc301167"
>>> twokenize.tokenizeRawTweetText(input)
['Clowns', 'are', 'pretty', 'gross', 'tho', 'O.o', '(', '"I'm", 'afraid', 'of', 'clowns', ':p', ')', 'ask.fm/a/cc301167']
```



← input


← output

Tool: twokenize.py

 **GitHub, Inc. [US]** <https://github.com/myleott/ark-twokenize-py/blob/master/twokenize.py>

 This repository Pull requests Issues Gist

 **myleott / ark-twokenize-py** 

 branch: **master** **ark-twokenize-py / twokenize.py**

 **myleott** on Apr 29, 2013 Initial commit

1 contributor

Executable File | 300 lines (247 sloc) | 12.993 kB

[Raw](#) [Edit](#)

```
1 # -*- coding: utf-8 -*-
2 """
3 Twokenize -- a tokenizer designed for Twitter text in English and some other European languages.
```


Tool: twokenize.py

```
3 Twokenize -- a tokenizer designed for Twitter text in English and some other European languages.
4 This tokenizer code has gone through a long history:
5
6 (1) Brendan O'Connor wrote original version in Python, http://github.com/brendano/tweetmotif
7     TweetMotif: Exploratory Search and Topic Summarization for Twitter.
8     Brendan O'Connor, Michel Krieger, and David Ahn.
9     ICWSM-2010 (demo track), http://brenocon.com/oconnor\_krieger\_ahn.icwsm2010.tweetmotif.pdf
10 (2a) Kevin Gimpel and Daniel Mills modified it for POS tagging for the CMU ARK Twitter POS Tagger
11 (2b) Jason Baldridge and David Snyder ported it to Scala
12 (3) Brendan bugfixed the Scala port and merged with POS-specific changes
13     for the CMU ARK Twitter POS Tagger
14 (4) Tobi Owoputi ported it back to Java and added many improvements (2012-06)
15
16 Current home is http://github.com/brendano/ark-tweet-nlp and http://www.ark.cs.cmu.edu/TweetNLP
```

Tokenization

- main techniques:
 - hand-crafted rules as regular expressions

Regular Expression

- a pattern matching language
- invented by American Mathematician Stephen Kleene in the 1950s
- used for search, find, replace, validation ... (very frequently used when dealing with strings)
- supported by most programming languages
- easy to learn, but hard to master

Regular Expression

| | |
|-----|---------------------------|
| 147 | Hashtag = "[a-zA-Z0-9_]+" |
|-----|---------------------------|

- [] indicates a set of characters:
 - [amk] will match 'a', 'm', or 'k'
 - [a-z] will match any lowercase letter ('abcdefghijklmnopqrstuvwxyz')
 - [a-zA-Z0-9_] will match any letter or digit or '_'
- + matches 1 or more repetitions of preceding RE

Regular Expression

| | |
|-----|---------------------------|
| 147 | Hashtag = "[a-zA-Z0-9_]+" |
|-----|---------------------------|

- will match strings that:
 - start with a '#'
 - follow with one or more letters/digits/'_'

Regular Expression

| | |
|-----|---------------------------|
| 147 | Hashtag = "[a-zA-Z0-9_]+" |
|-----|---------------------------|

```
>>> import re
>>> Hashtag = "[a-zA-Z0-9_]+"
>>> hashtagpattern = re.compile(Hashtag)
>>> hashtagpattern.findall("So that's what #StarWars")
['#StarWars']
```

Regular Expression

| | |
|-----|------------------------|
| 133 | Hearts = "(?:<+/?3+)+" |
|-----|------------------------|

- ‘+’ matches 1 or more repetitions of the preceding RE
 - ‘<+’ matches ‘<’, ‘<<’, ‘<<<’ ...
 - ‘3+’ matches ‘3’, ‘33’, ‘333’ ...
- ‘?’ matches 0 or 1 repetitions of the preceding RE
 - ‘/?’ matches ‘/’ or nothing (so handles ‘</3’)
- (?: ...) is a non-capturing version of (...)
- (...) matches whatever RE is inside the parentheses

Regular Expression

| | |
|-----|------------------------|
| 133 | Hearts = "(?:<+/?3+)+" |
|-----|------------------------|

- will match strings that:
 - start with one or more '<'
 - then maybe a '/'
 - then one or more '3'
 - and maybe repetitions of the above

Regular Expression

| | |
|-----|------------------------|
| 133 | Hearts = "(?:<+/?3+)+" |
|-----|------------------------|

```
>>> import re
>>> Hearts = "(?:<+/?3+)+"
>>> heartspattern = re.compile(Hearts)
>>> heartspattern.findall("I <3 u <3<333333")
['<3', '<3<333333']
>>> heartspattern.findall("sooo sad </3")
['</3']
```

Regular Expression

- learn more (<https://docs.python.org/2/library/re.html>)

Python » 2.7.10 Documentation » The Python Standard Library » 7. String Services »

previous | next | modules | index

Table Of Contents

- 7.2. `re` — Regular expression operations
 - 7.2.1. Regular Expression Syntax
 - 7.2.2. Module Contents
 - 7.2.3. Regular Expression Objects
 - 7.2.4. Match Objects
 - 7.2.5. Examples
 - 7.2.5.1. Checking For a Pair
 - 7.2.5.2. Simulating `scanf()`
 - 7.2.5.3. `search()` vs. `match()`
 - 7.2.5.4. Making a Phonebook
 - 7.2.5.5. Text Munging
 - 7.2.5.6. Finding all Adverbs
 - 7.2.5.7. Finding all Adverbs and their Positions
 - 7.2.5.8. Raw String Notation

7.2. `re` — Regular expression operations

This module provides regular expression matching operations similar to those found in Perl. Both patterns and strings to be searched can be Unicode strings as well as 8-bit strings.

Regular expressions use the backslash character (`'\'`) to indicate special forms or to allow special characters to be used without invoking their special meaning. This collides with Python's usage of the same character for the same purpose in string literals; for example, to match a literal backslash, one might have to write `'\\'` as the pattern string, because the regular expression must be `\\`, and each backslash must be expressed as `\\` inside a regular Python string literal.

The solution is to use Python's raw string notation for regular expression patterns; backslashes are not handled in any special way in a string literal prefixed with `'r'`. So `r"\n"` is a two-character string containing `'\'` and `'n'`, while `"\n"` is a one-character string containing a newline. Usually patterns will be expressed in Python code using this raw string notation.

It is important to note that most regular expression operations are available as module-level functions and `RegexObject` methods. The functions are shortcuts that don't require you to compile a regex object first, but miss some fine-tuning parameters.

7.2.1. Regular Expression Syntax

Tokenization

- for Twitter, additionally need to handle:
 - emoticons, urls, #hashtags, @mentions ...

```
>>> import twokenize
>>> input = "Clowns are pretty gross tho O.o (I'm afraid of clowns :p) ask.fm/a/cc301167"
>>> twokenize.tokenizeRawTweetText(input)
['Clowns', 'are', 'pretty', 'gross', 'tho', 'O.o', '(', '"I'm", 'afraid', 'of', 'clowns', ':p', ')', 'ask.fm/a/cc301167']
```

← input


← output

Tokenization


- language dependent

| | |
|-----------------|---|
| 下雨天留客天留我不留 | Unpunctuated Chinese sentence |
| 下雨、天留客。天留、我不留！ | <i>It is raining, the god would like the guest to stay. Although the god wants you to stay, I do not!</i> |
| 下雨天、留客天。留我不？ 留！ | <i>The rainy day, the staying day. Would you like me to stay? Sure!</i> |
| 我喜欢新西兰花 | Unsegmented Chinese sentence |
| 我 喜欢 新西兰 花 | <i>I like New Zealand flowers</i> |
| 我 喜欢 新 西兰花 | <i>I like fresh broccoli</i> |

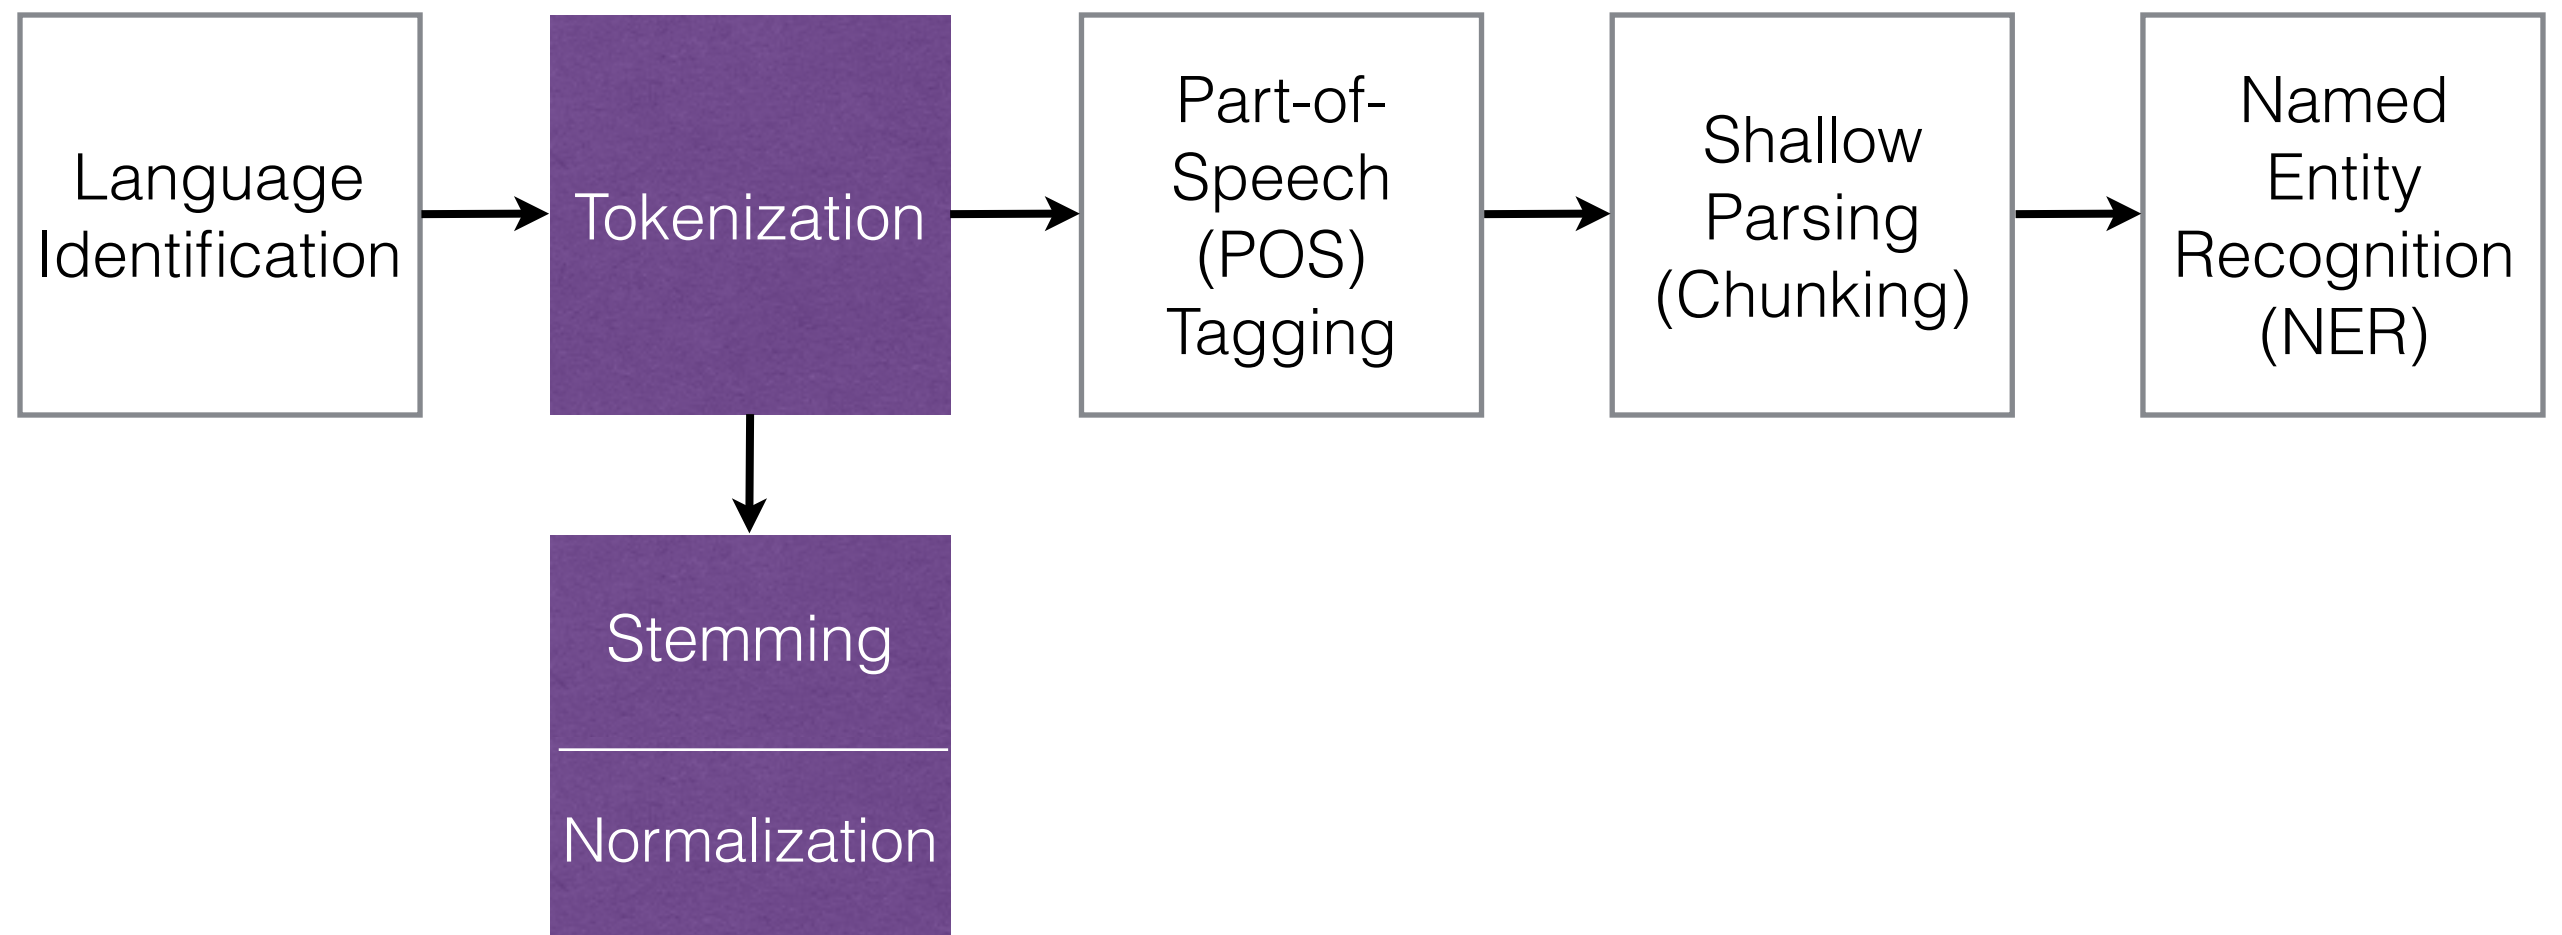
NLP Pipeline



NLP Pipeline



NLP Pipeline



Stemming

- reduce inflected words to their word stem, base or root form (not necessarily the morphological root)
- studied since the 1960s

```
>>> from nltk.stem.porter import PorterStemmer
>>> porter_stemmer = PorterStemmer()
>>> porter_stemmer.stem('maximum')
'maximum'
>>> porter_stemmer.stem('presumably')
'presum'
>>> porter_stemmer.stem('multiply')
'multipli'
```



Stemming

- different steamers: Porter, Snowball, Lancaster ...
- WordNet's built-in lemmatized (dictionary-based)

```
>>> from nltk.stem import WordNetLemmatizer
>>> wordnet_lemmatizer = WordNetLemmatizer()
>>> wordnet_lemmatizer.lemmatize('leaves', pos='n')
'leaf'
>>> wordnet_lemmatizer.lemmatize('leaves', pos='v')
'leave'
```

Stemming

- language dependent



Text Normalization

- convert non-standard words to standard

Original tweet

@USER, r u cuming 2 MidCorner dis Sunday?

Normalized tweet

@USER, are you coming to MidCorner this Sunday?

Original tweet

Still have to get up early 2mr thou 😞 so Gn 😴

Normalized tweet

Still have to get up early tomorrow though 😞 so Good night 😴

Source: Baldwin, de Marneffe, Han, Kim, Ritter, Xu
Shared Tasks of the 2015 Workshop on Noisy User-generated Text:
Twitter Lexical Normalization and Named Entity Recognition

Text Normalization

- types of non-standard words in 449 English tweets:

| Category | Ratio | Example |
|---------------------|--------|----------------------|
| letter&number | 2.36% | b4 → before |
| letter | 72.44% | shuld → should |
| number substitution | 2.76% | 4 → for |
| slang | 12.20 | lol → laugh out loud |
| other | 10.24% | sucha → such a |

most non-standard words are morphophonemic “errors”

Source: Bo Han and Timothy Baldwin

“Lexical normalisation of short text messages: Makn sens a #twitter” ACL 2011

A Normalization Lexicon

- automatically derived from Twitter data + dictionary

| | | |
|-------|------------|-----------|
| 41169 | costumess | costumes |
| 41170 | nywhere | anywhere |
| 41171 | sandwich | sandwich |
| 41172 | aleksander | alexander |
| 41173 | juns | jun |
| 41174 | showi | showing |
| 41175 | washinq | washing |
| 41176 | jscript | script |
| 41177 | fundin | funding |
| 41178 | itxted | fitted |
| 41179 | cheeeap | cheap |
| 41180 | fawesome | awesome |
| 41181 | untalented | talented |
| 41182 | | |


Performance

Precision = 0.847

Recall = 0.630

F1-Score = 0.723

F-measure



How many selected items are relevant?


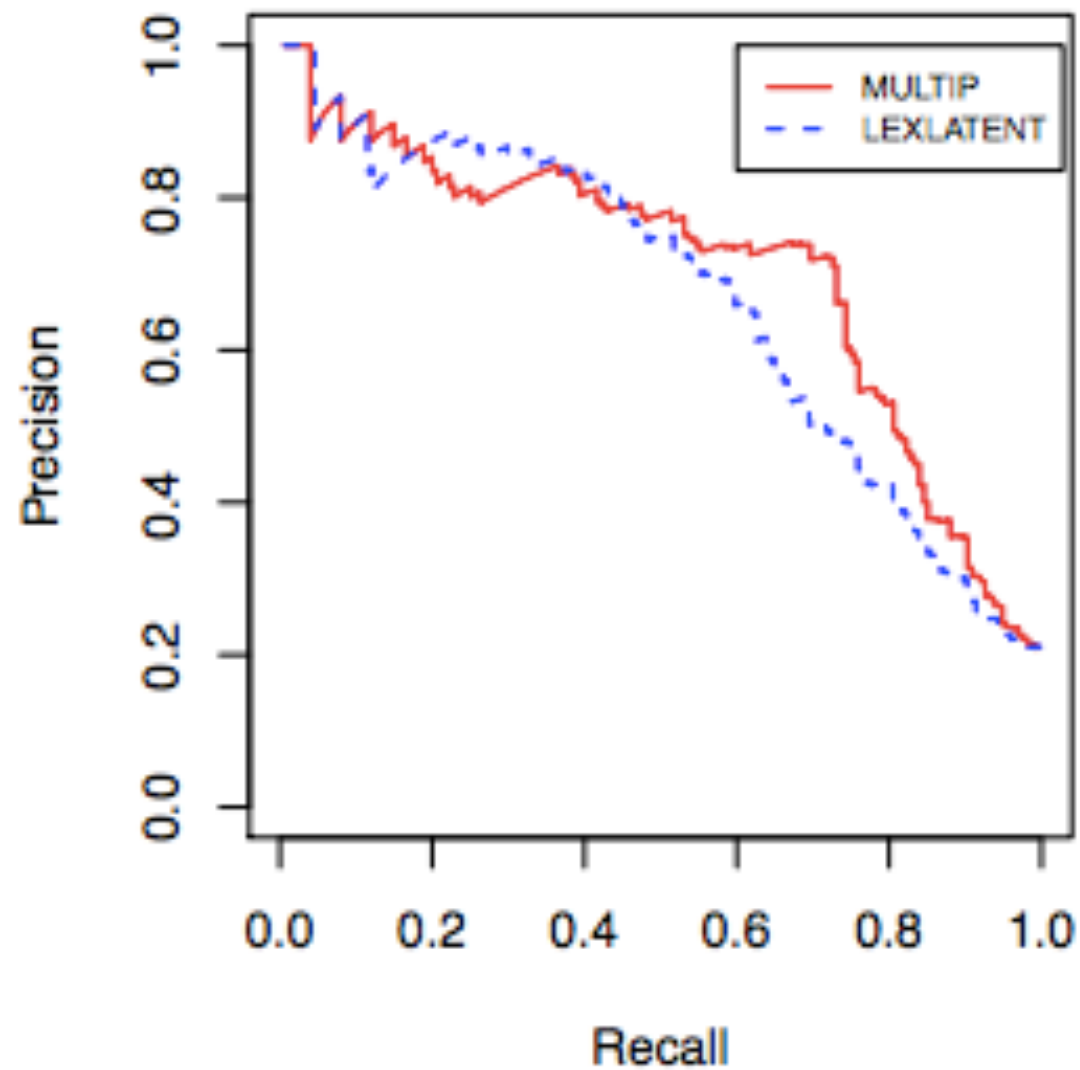
$$\text{Precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$

$$F_1 = \frac{2 \cdot \text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

Precision-Recall Curve



showing the trade-off, no threshold picking

Phrase-level Normalization

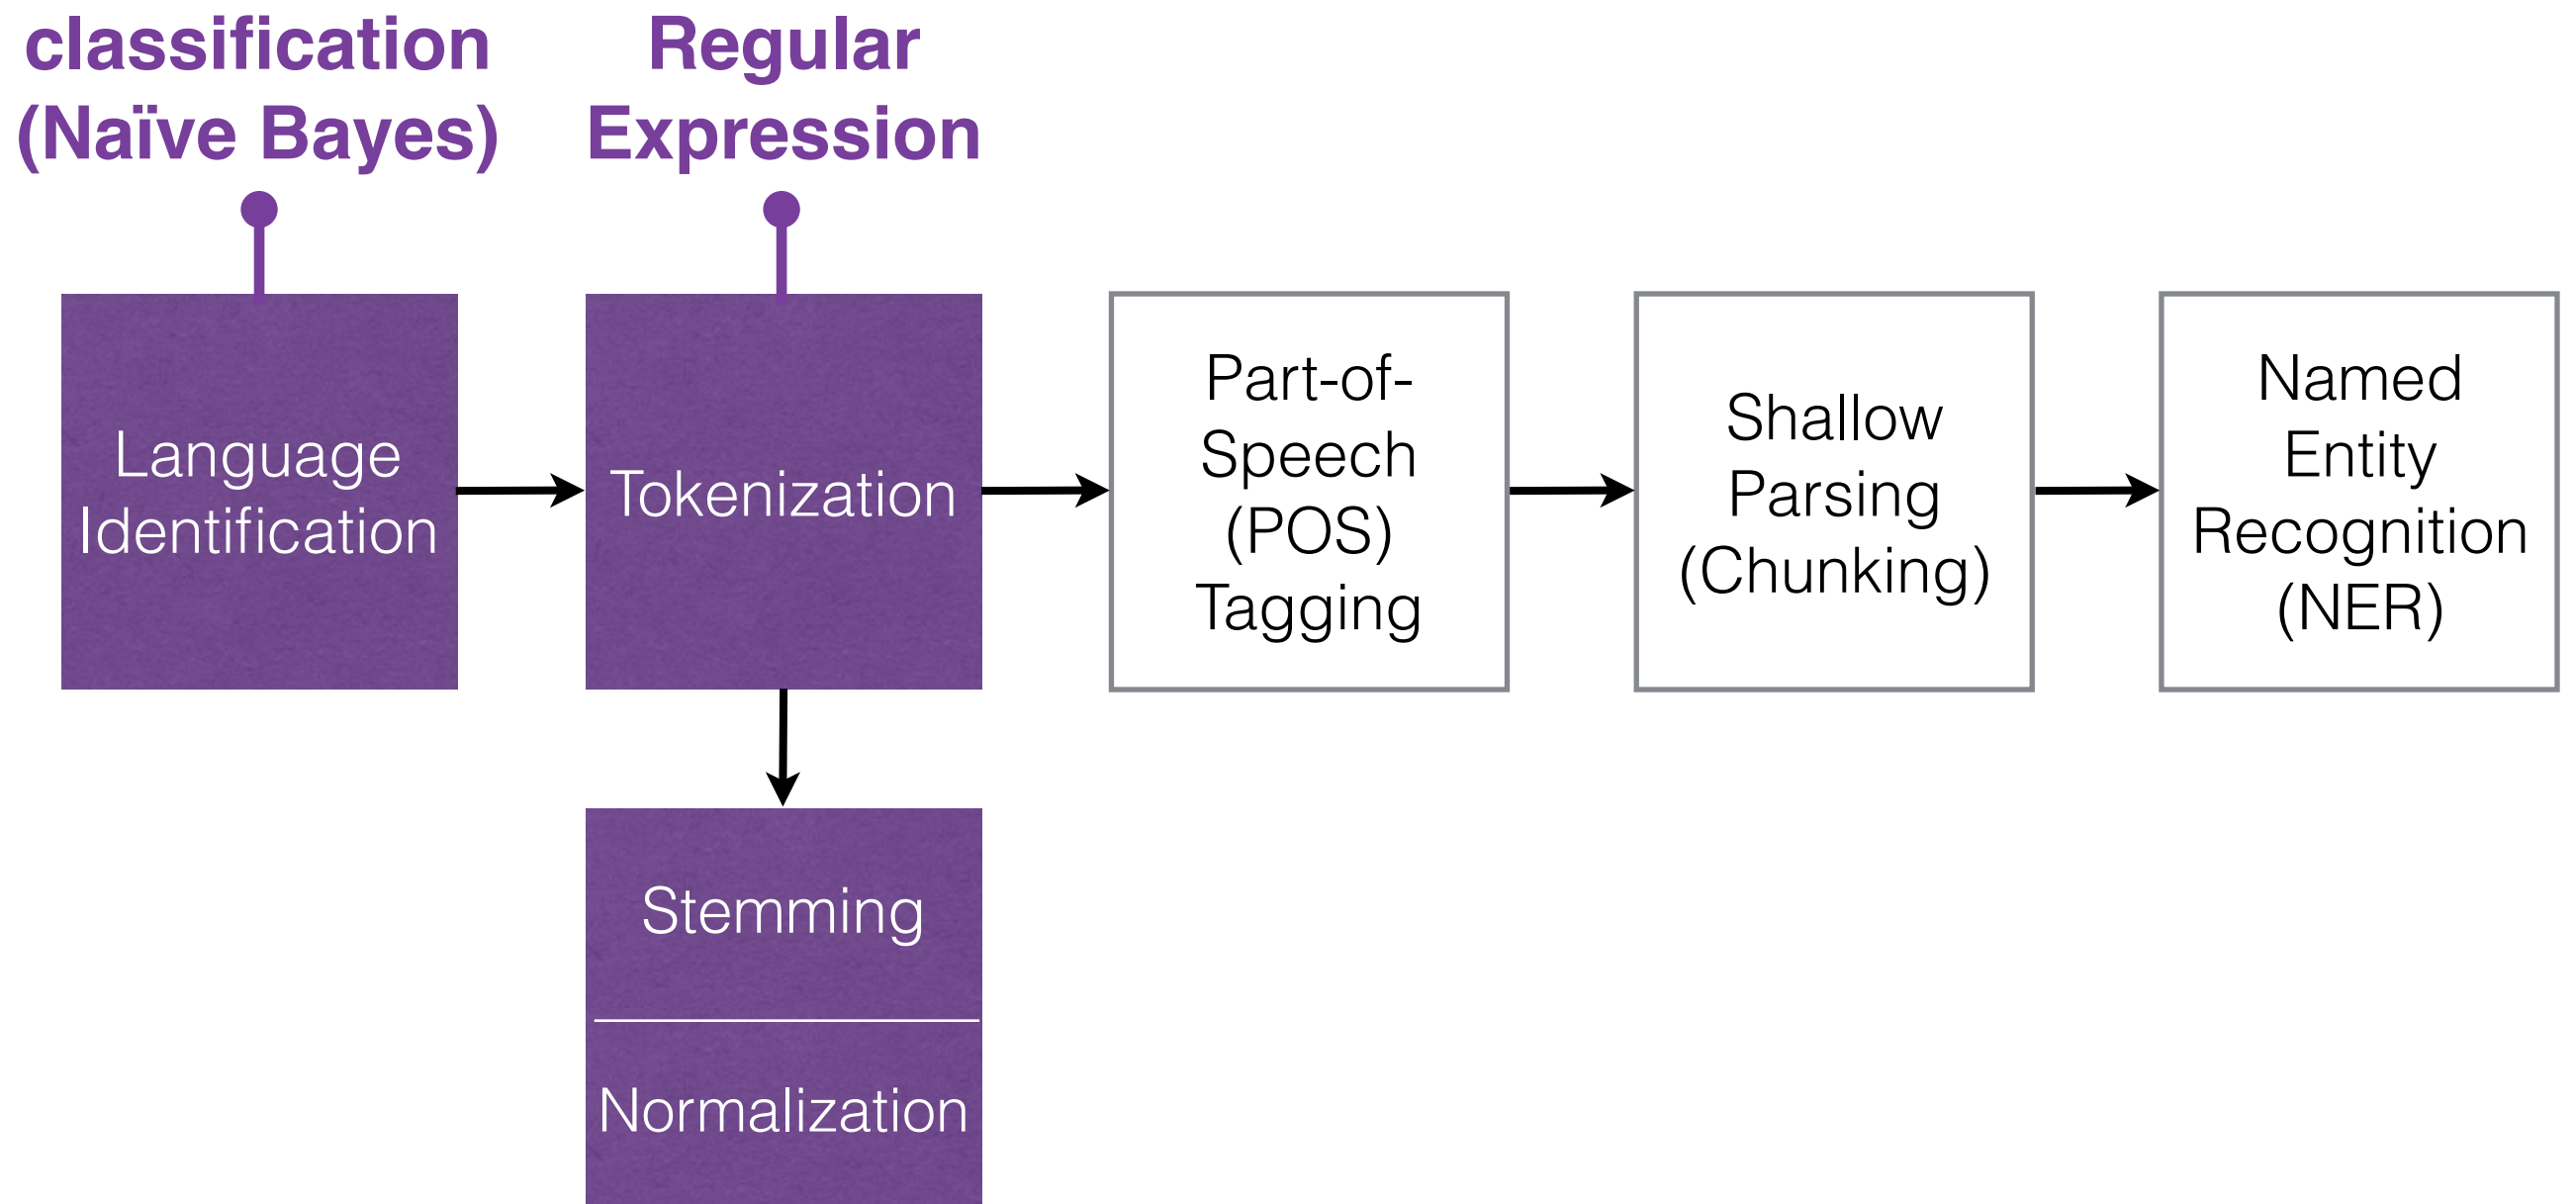
- word-level normalization is insufficient for many cases:

in-vocabulary words

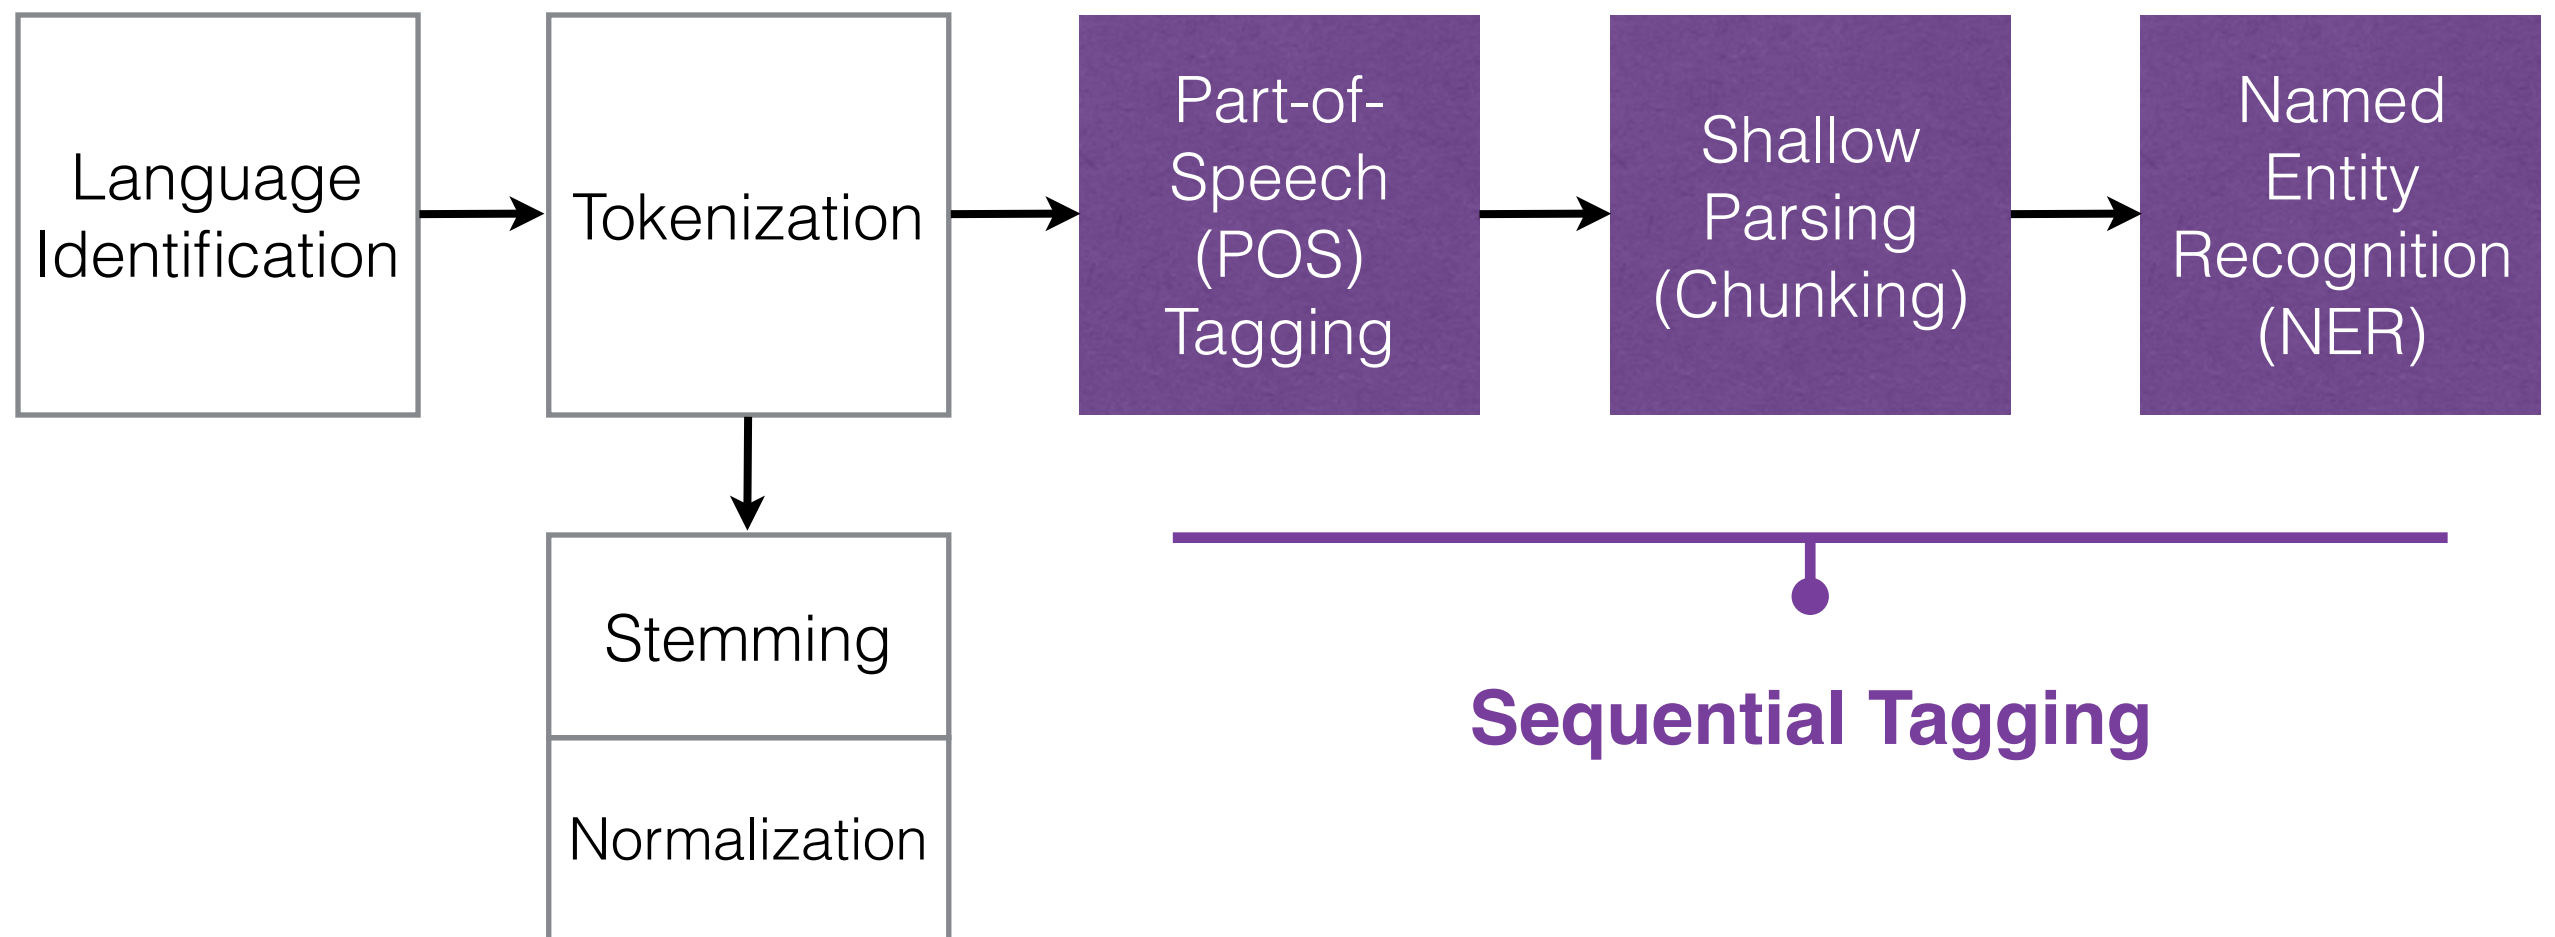
| Category | Example |
|---------------|--|
| 1-to-many | everytime → every time |
| incorrect IVs | can't want for → can't wait for |
| grammar | I'm going a movie → I'm going to a movie |
| ambiguities | 4 → 4 / 4th / for / four |

Source: Wei Xu, Alan Ritter, Ralph Grishman


Summary



Next Lecture



Thank You!



Instructor: Wei Xu

www.cis.upenn.edu/~xwe/

Course Website: socialmedia-class.org